

Job Title: AI Intern

Location: Bengaluru

Company: Ambitio

Company Description:

Ambitio is building an AI platform for higher education across the globe. Backed by First Cheque, WFC & some of the best angels, we empower people to make informed decisions and improve their odds of getting into top institutions. Our users have received admits from NYU, U Penn, Cornell University, Columbia University, Carnegie Mellon University, UC Berkeley, ICL, Warwick Business School, and other top QS universities.

Are you passionate about education and helping students achieve their dreams of studying abroad? Join our dynamic team as a Junior Admission Expert and help us guide aspiring students to their next adventure!

About Ambitio: Copilot for Education Abroad

At Ambitio, we're building India's first AI-powered global admissions platform, designed to make one of life's most complex decisions simpler, smarter, and more accessible.

Our mission is to democratize access to top-tier global education by replacing guesswork, bias, and expensive consulting with data-driven, personalized AI guidance.

Founded by IIT BHU alumni and backed by First Cheque, BLinC Invest, and We Founder Circle, we are a product-led company where AI is not a feature, it's the core.

- The Brain: Our AI engine is trained on 10M+ data points and 1M+ successful applications, delivering predictive Admit Probability scores
- The Ecosystem: From LLM-powered Essay Optimizers to advanced RAG pipelines and recommendation systems
- The Impact: 40,000+ students guided to admits from Stanford, Oxford, NYU, and the Ivy League
- We're now scaling our AI systems to support the next million users.

About the Role: AI Engineering Intern

Most AI internships stop at prompt engineering. This one doesn't.

We're looking for AI Engineering Interns who want to move beyond notebooks and experiments and work on real, production-grade AI systems used by thousands of users. This is not a "learn and observe" role. This is a build, optimize, and deploy role. You'll function as a high-ownership AI engineer, contributing directly to the core intelligence layer of the product.

What you'll work on

You'll work on real-world AI problems that go far beyond basic LLM usage:

- **Production AI Systems** : Build, optimize, and deploy models that serve real users (not just local notebooks)
- **RAG Pipelines & Retrieval Systems** : Design high-accuracy, low-latency retrieval systems for complex decision-making
- **Agentic Workflows** : Build autonomous AI systems capable of reasoning over messy, unstructured data
- **Model Optimization** : Experiment with fine-tuning techniques (PEFT, prompt tuning, etc.) to improve output quality and reliability
- **Predictive Modeling** : Improve the algorithms behind Admit Probability for highly competitive universities

Why this role matters :

This is one of those rare opportunities where:

- **The impact is real** → Your models directly influence life-changing decisions
- **The system is broken** → Admissions today is chaotic, biased, and opaque
- **The opportunity is massive** → We're building the intelligence layer for a \$10B+ problem

Every model you improve:

- helps a student make a better decision
- reduces uncertainty in a high-stakes journey
- and unlocks opportunities they didn't know existed

If you want to build AI that actually ships and matters, this is it.

What makes you a great fit?

Must-haves:

- Strong ML/DL fundamentals (you understand how models work, not just APIs)
- Solid grasp of DSA, probability, and system thinking
- Proficiency in Python
- Hands-on experience with PyTorch / TensorFlow / Scikit-learn

Bonus:

- Experience with RAG systems, LangChain, or LLM pipelines
- Familiarity with vector databases (Pinecone, Milvus, etc.)
- Kaggle, research work, or open-source contributions

Mindset:

- You're a builder, you learn by doing, not waiting
- You care about shipping working systems, not just experiments
- You enjoy debugging complex issues and figuring things out
- You take ownership of outcomes, not just tasks

Don't apply if

- You need constant guidance and a highly structured environment
- You prefer theoretical research over shipping real systems
- You're looking for a "safe" internship primarily for brand value

How you'll contribute

- Own AI features end-to-end: from experimentation → deployment
- Ship production-grade models used by thousands of users
- Build and improve intelligent systems (RAG, agents, recommendations)
- Debug and optimize models in real-world conditions
- Work closely with founders and core engineers
- Continuously improve performance, accuracy, and reliability

What your day will look like

No isolated experiments. No toy problems. A typical day might include:

- Turning ambiguous AI problems into structured solutions
- Building and deploying models into production
- Debugging model failures or hallucinations
- Improving retrieval quality or response accuracy
- Collaborating with engineers and founders on product decisions
- Shipping fast, testing, and iterating

How to apply :

Share your resume with us. Show us what you've built.

- Share your GitHub / ML project / deployed system
- Walk us through your approach and decisions
- Help us understand how you think

Hiring Process

We keep it fast, focused, and meaningful:

- **R1:** Intro Call → Background, projects, and curiosity
- **R2:** Tech Round → ML fundamentals + problem-solving + system thinking
- **R3:** Founder Round → Ownership, ambition, and builder mindset

Offer

If there's a strong match:

- You'll receive a quick decision and fast offer rollout
- You'll join as a high-ownership AI intern (not a support role)
- Exceptional performance can lead to a full-time AI Engineer offer

Employment Type: Internship

Working Days -6 (Mon- Sat)

Timings - 11AM-8PM

Location - HSR, Bengaluru